

Jay Singh

Jays.iitkgp@gmail.com | +91-8651274328 | [linkedin.com/in/jay-singh-ds/](https://www.linkedin.com/in/jay-singh-ds/) | github.com/JaySingh79

EDUCATION

IIT Kharagpur

Kharagpur, India

Bachelor of Technology (B.Tech)

July 2026

- **Relevant Coursework:** Probability and Statistics, Machine Learning Foundations, Applied Mathematics, LDA.

WORK EXPERIENCE

Pibit AI (YC W21)

Jan – May 2026

AI Intern (Research)

Bengaluru, India

- Led **Template Mapper**, the classification core of a **loss-run** API system with **>91%** accuracy, over **850** labels
- **Productionized** the end-to-end extraction, ingestion, preprocessing, prompt optimization, evaluation, and QC
- Launched Document Handler with **OCR** framework for ingestion failure detection, saving **4+** operational days

Pascal AI Labs

Nov – Dec 2025

ML Intern

Remote

- Pioneered DFR-Tracer to interpret restatements in SEC 10-K filings, with **LangGraph** achieving 90% accuracy
- Optimized multi-hop **lineage graph**, with Pgvector and fuzzy alignment enabling high-precision RAG retrieval
- Eliminated black-box LLM outputs by **hybrid reranking** with source citation and hallucination rejection loops

Innovaccer

Jun – Aug 2024

Data Science Intern

Noida, India

- Built a classification model identifying **quality measures/risk modules** from client queries with **85%** accuracy
- Enhanced LangChain-Milvus RAG system over 150 pages, refining chunking and cutting latency by **73%** to 8.1s
- Streamlined a **multi-chain** pipeline routing queries across **5+** classifiers in parallel at **7s** average latency at scale

Cambridge Judge Business School

May – Sep 2023

Research Apprenticeship

Remote

- Scaled a **DAG-based agentic** workflow, cutting manual research by **90%** to generate data-driven business ideas
- Designed a shared **agent memory** layer for reasoning tasks cutting tool calls by 44% versus no-memory version
- Refined agent execution through parallel task scheduling and response caching, slicing inference cost and latency

PROJECTS

Attention Based Federated Learning Approach for IDS | CSE IIT Kharagpur

May 2025

- Achieved **80%** zero-shot detection accuracy with a 4-layer transformer classifier, residual MLP, and LayerNorm
- Enriched NSL-KDD dataset with metadata-rich **LLaMA** output for **60,000** entries and contextual vector labels
- Applied **FedAdam** optimizer, stabilizing convergence rates by 25% and improving generalization against attacks

ConvoHire AI | AI Powered Recruitment Model

Feb 2025

- Designed a 6-stage **LLM interview** pipeline - scoring, triggering adaptive follow-ups, tracking topic coverage
- Coded hybrid exact + vector JD-resume skill matching (10+ page support), driving targeted question generation
- Added an **LLM-as-judge** weighing candidate's strength, tech-stack relevance, and domain expertise for verdict

"Attention is All You Need" Paper Replica | Built From Scratch

Jan 2025

- Crafted an **encoder-decoder** Transformer using PyTorch with a 10k vocab size and **BPE-based tokenization**
- Utilized multi-head attention with parallel Q/K/V projections to reproduce the original Transformer system
- Implemented position-wise **FFN** with ReLU activations, **residual connections**, **LayerNorm** on each sublayer

CERTIFICATIONS & SKILLS

- **Certifications:** PyTorch for Deep learning, Career Essentials on Gen AI, Advanced Learning Algorithms
- **Technologies/Frameworks:** Python, TensorFlow, PyTorch, AWS, Snowflake, FastAPI, LangChain, Django, Docker, SQL, RAG, OpenTelemetry, MLflow, ReAct, Git, GitHub, Hugging Face, DSPy, XGBoost, LightGBM
- **Skills:** Applied ML, Recommender Systems, Statistics, Agentic AI, RLHF, Fine Tuning, LLM Evaluation, OCR, CNN, NLP, Federated Learning, Document Extraction, Feature Engineering, MLOps, Transfer Learning